


Computer Architecture

Trần Trọng Hiếu

Information Systems Department
Faculty of Information Technology

VNU - UET

hieutt@vnu.edu.vn






Enhancing Performance with Pipelining


Pipelining



- Start work ASAP!! Do not waste time!

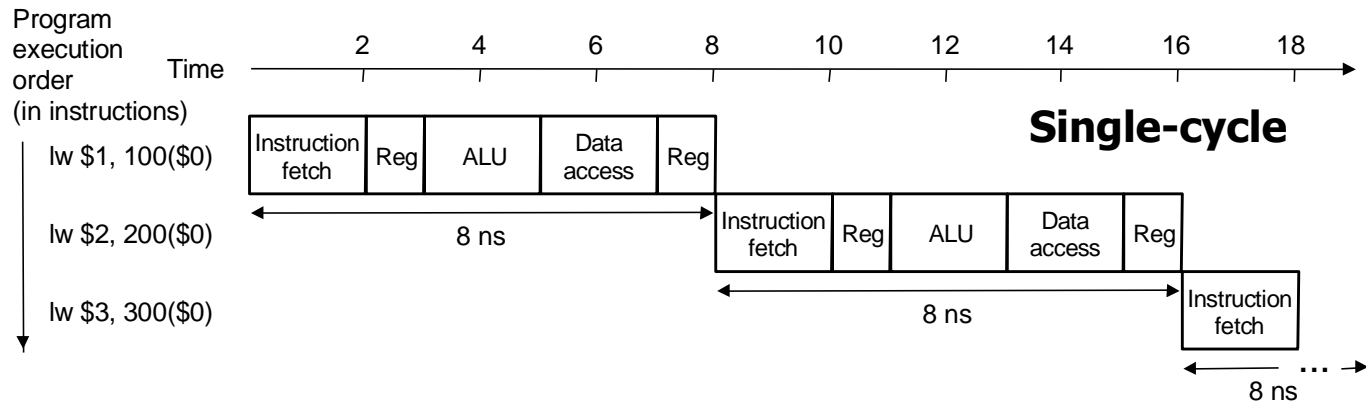


Assume 30 min. each task – wash, dry, fold, store – and that separate tasks use separate hardware and so can be overlapped

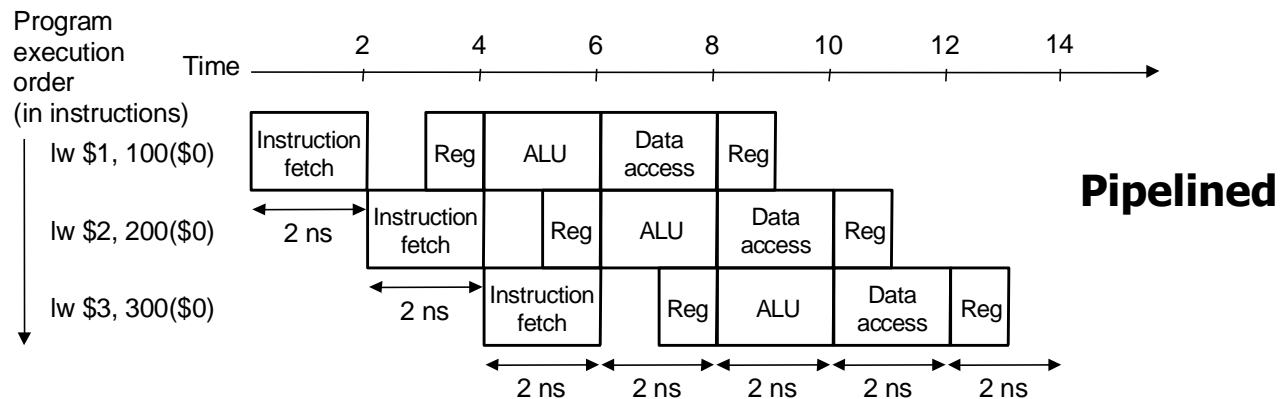




Pipelined vs. Single-Cycle Instruction Execution: the Plan



Assume 2 ns for memory access, ALU operation; 1 ns for register access: therefore, single cycle clock 8 ns; pipelined clock cycle 2 ns.





Pipelining: Keep in Mind

- Pipelining *does not reduce latency* of a single task, it *increases throughput* of entire workload
- Pipeline rate *limited by longest stage*
 - *potential* speedup = number pipe stages
 - *unbalanced lengths* of pipe stages reduces speedup
- Time to *fill* pipeline and time to *drain* it – when there is *slack* in the pipeline – reduces speedup



Example Problem

- *Problem: for the laundry fill in the following table when*
 1. *the stage lengths are 30, 30, 30 30 min., resp.*
 2. *the stage lengths are 20, 20, 60, 20 min., resp.*

Person	Unpipelined finish time	Pipeline 1 finish time	Ratio unpipelined to pipeline 1	Pipeline 2 finish time	Ratio unpipelelined to pipeline 2
1					
2					
3					
4					
n					

- *Come up with a formula for pipeline speed-up!*

Pipelining MIPS




- What makes it easy with MIPS?
 - all *instructions are same length*
 - so fetch and decode stages are similar for all instructions
 - just a *few instruction formats*
 - simplifies instruction decode and makes it possible in one stage
 - *memory operands appear only in load/stores*
 - so memory access can be deferred to exactly one later stage
 - *operands are aligned in memory*
 - one data transfer instruction requires one memory access stage

Pipelining MIPS

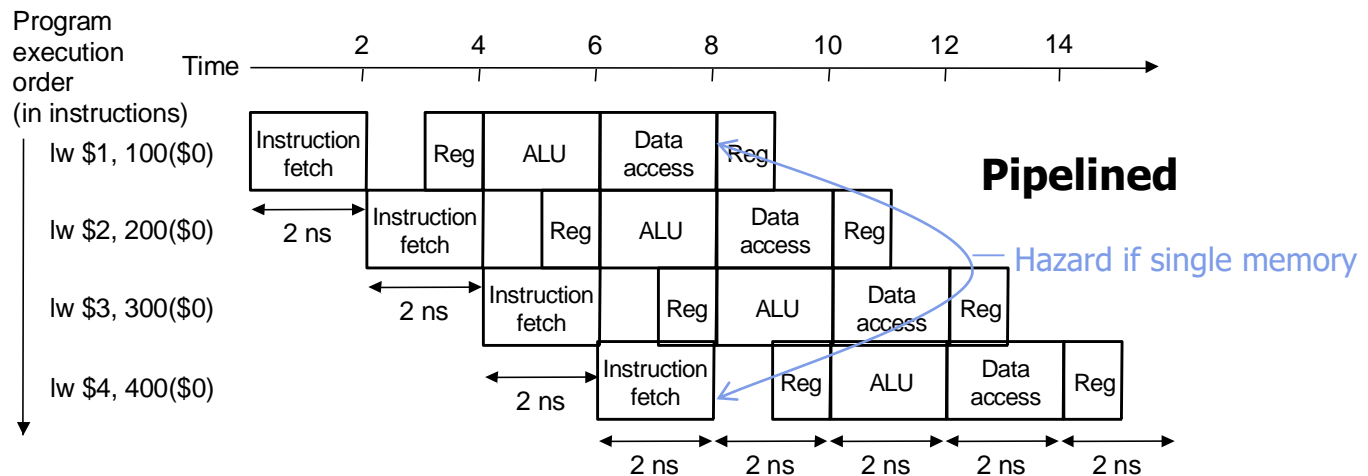


- What makes it hard?
 - *structural hazards*: different instructions, at different stages, in the pipeline want to use the same hardware resource
 - *control hazards*: succeeding instruction, to put into pipeline, depends on the outcome of a previous branch instruction, already in pipeline
 - *data hazards*: an instruction in the pipeline requires data to be computed by a previous instruction still in the pipeline
- Before actually building the pipelined datapath and control we first briefly examine these potential hazards individually...

Structural Hazards



- *Structural hazard*: inadequate hardware to simultaneously support all instructions in the pipeline in the same clock cycle
- E.g., suppose *single – not separate* – instruction and data memory in pipeline below with *one read port*
 - then a structural hazard between first and fourth lw instructions

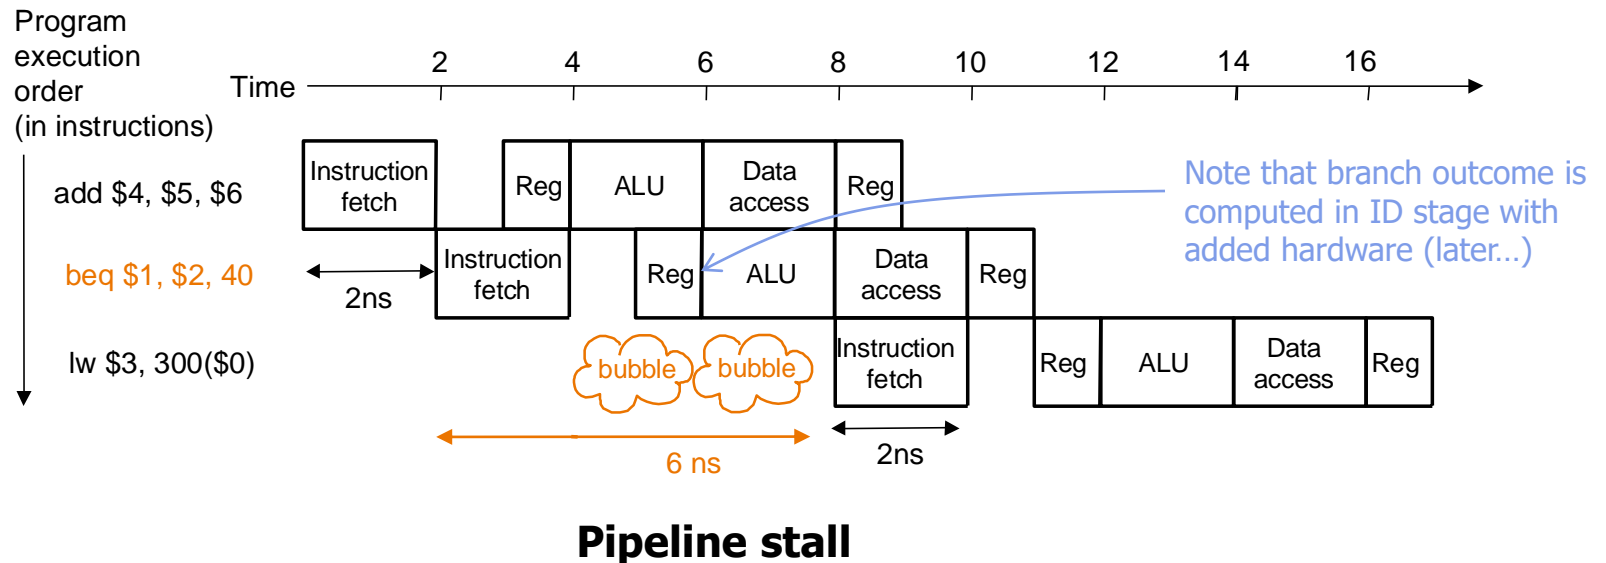


- *MIPS was designed to be pipelined*: structural hazards are easy to avoid!

Control Hazards



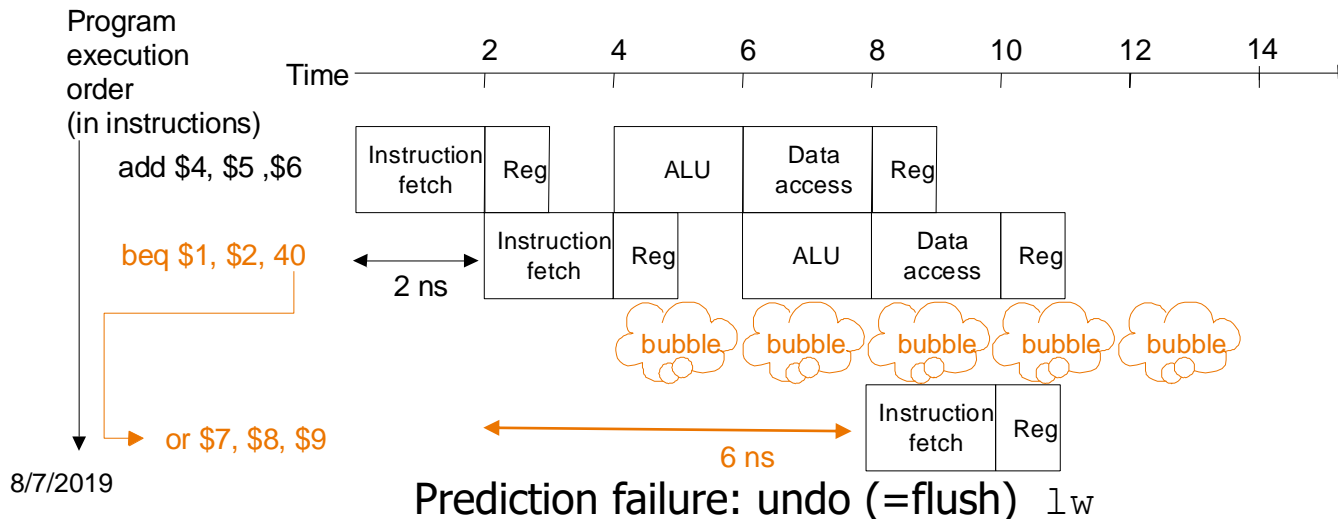
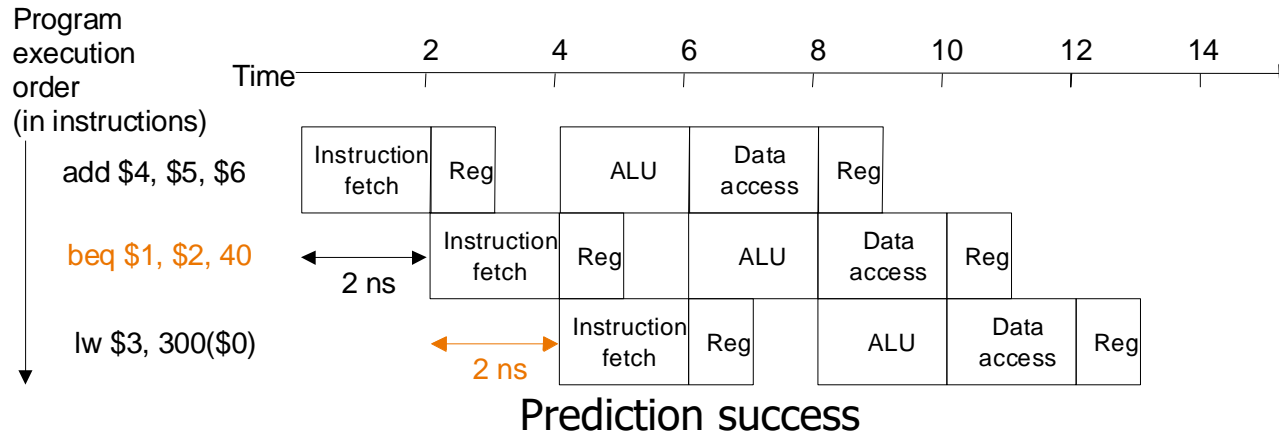
- *Control hazard*: need to make a decision based on the result of a previous instruction still executing in pipeline
- Solution 1 *Stall* the pipeline



Control Hazards



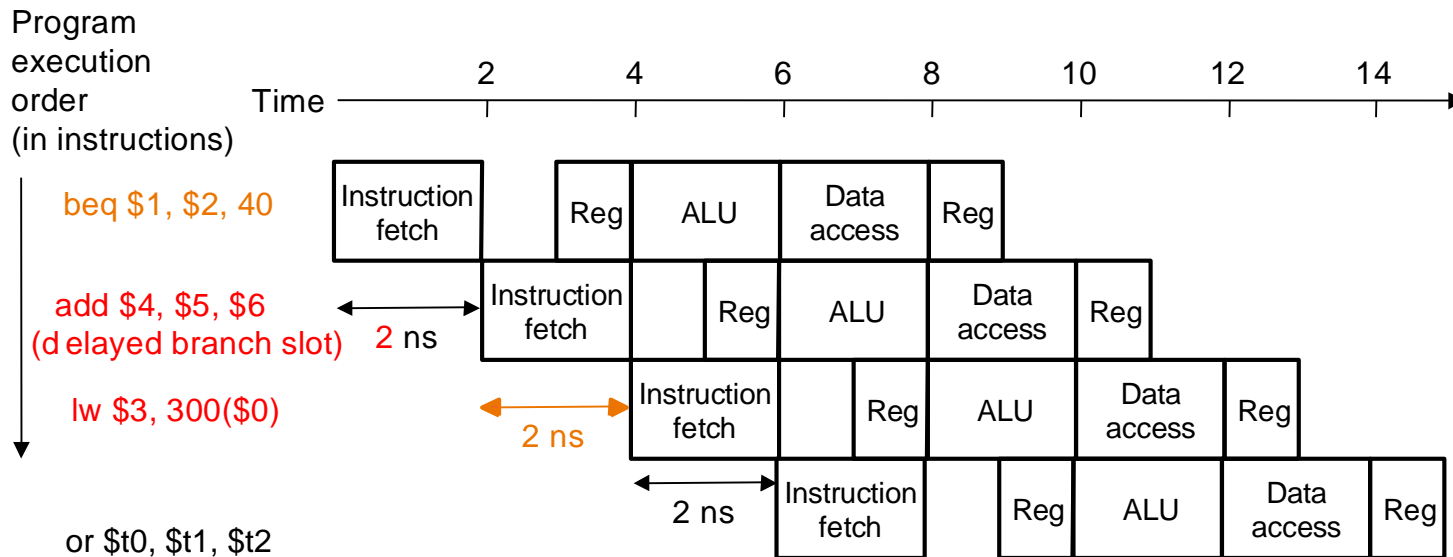
- Solution 2 *Predict* branch outcome
 - e.g., predict *branch-not-taken* :




Control Hazards




- Solution 3 *Delayed branch*: always execute the sequentially next statements with the branch executing after one instruction delay – compiler's job to find a statement that can be put in the slot that is independent of branch outcome
 - MIPS does this – but it is an option in SPIM (Simulator -> Settings)



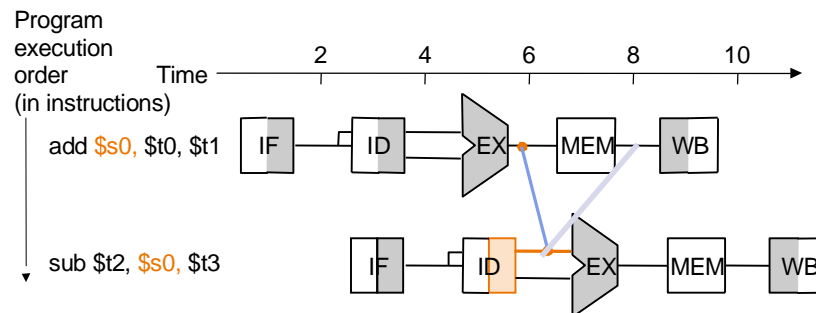
Data Hazards



- *Data hazard*: instruction needs data from the result of a previous instruction still executing in pipeline
- Solution *Forward* data if possible...



Instruction pipeline diagram:
shade indicates use –
left=write, right=read

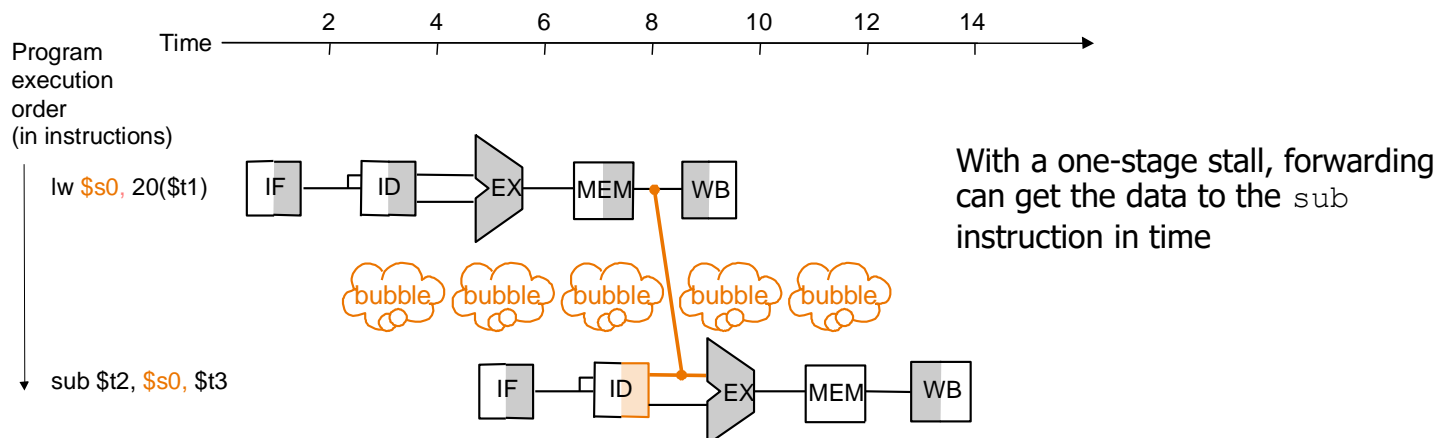
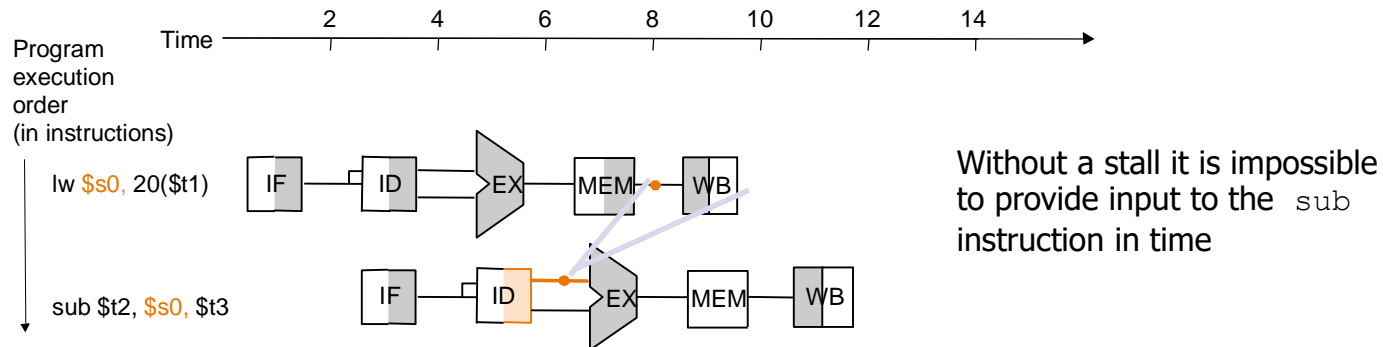


Without forwarding – blue line –
data has to go back in time;
with forwarding – red line –
data is available in time


Data Hazards



- Forwarding may not be enough
 - e.g., if an R-type instruction following a load uses the result of the load – called *load-use data hazard*



Reordering Code to Avoid Pipeline Stall (Software Solution)



- Example:

```
lw $t0, 0($t1)
lw $t2, 4($t1)
sw $t2, 0($t1)
sw $t0, 4($t1)
```

← Data hazard

- Reordered code:

```
lw $t0, 0($t1)
lw $t2, 4($t1)
sw $t0, 4($t1)
sw $t2, 0($t1)
```

← Interchanged

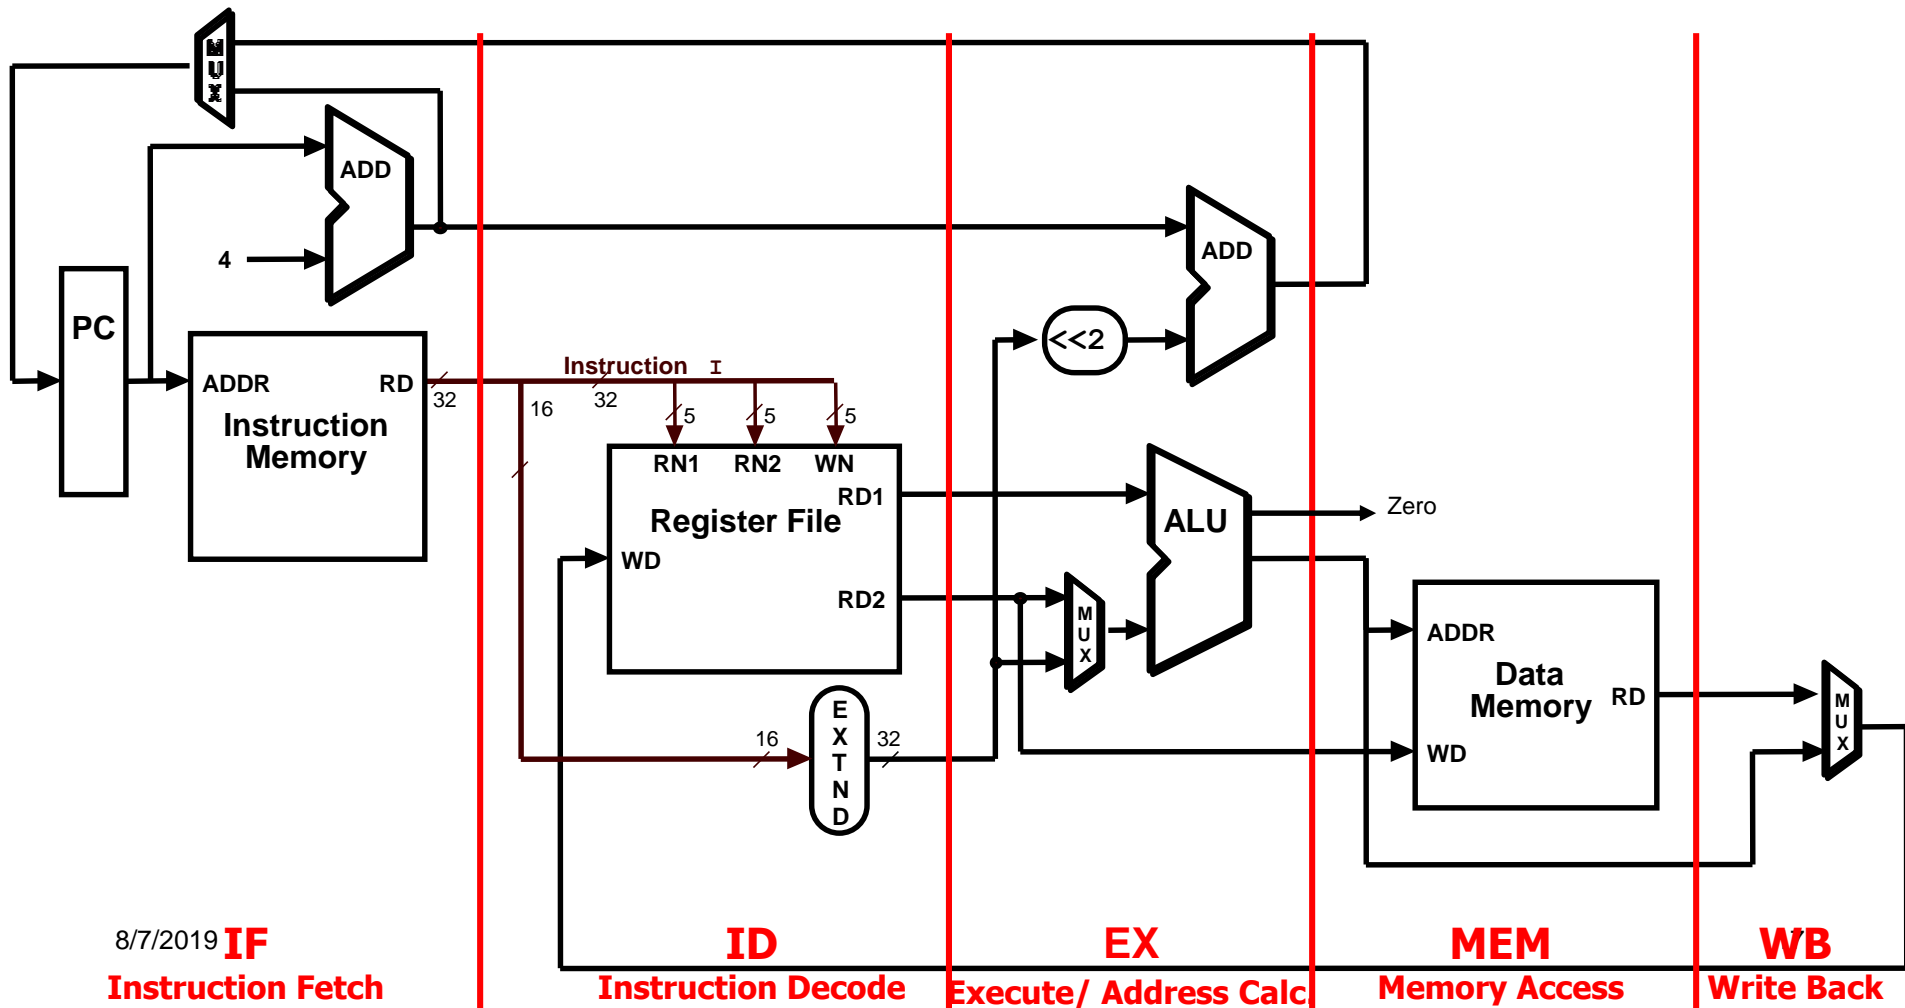
Pipelined Datapath



- We now move to actually building a pipelined datapath
- First recall the 5 steps in instruction execution
 1. Instruction Fetch & PC Increment (IF)
 2. Instruction Decode and Register Read (ID)
 3. Execution or calculate address (EX)
 4. Memory access (MEM)
 5. Write result into register (WB)
- Review: single-cycle processor
 - all 5 steps done in a single clock cycle
 - dedicated hardware required for each step
- *What happens if we break the execution into multiple cycles, but keep the extra hardware?*

Review - Single-Cycle Datapath

“Steps”

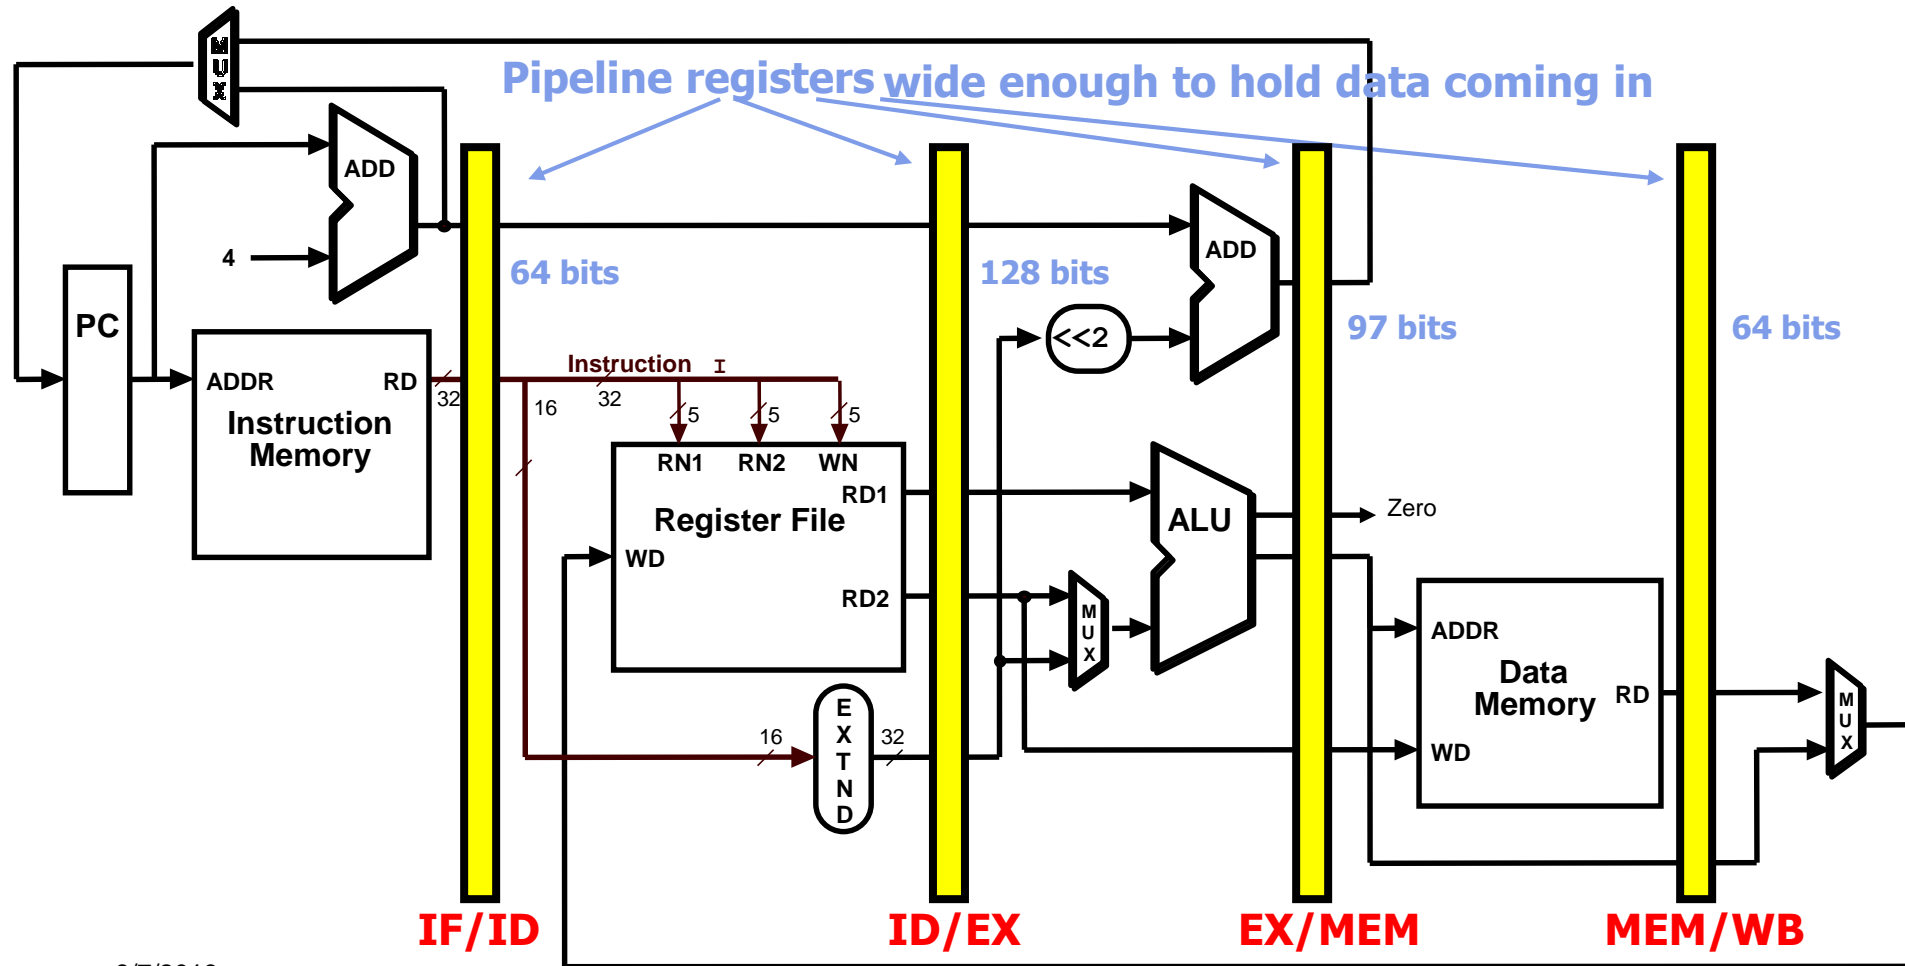


Pipelined Datapath – Key Idea

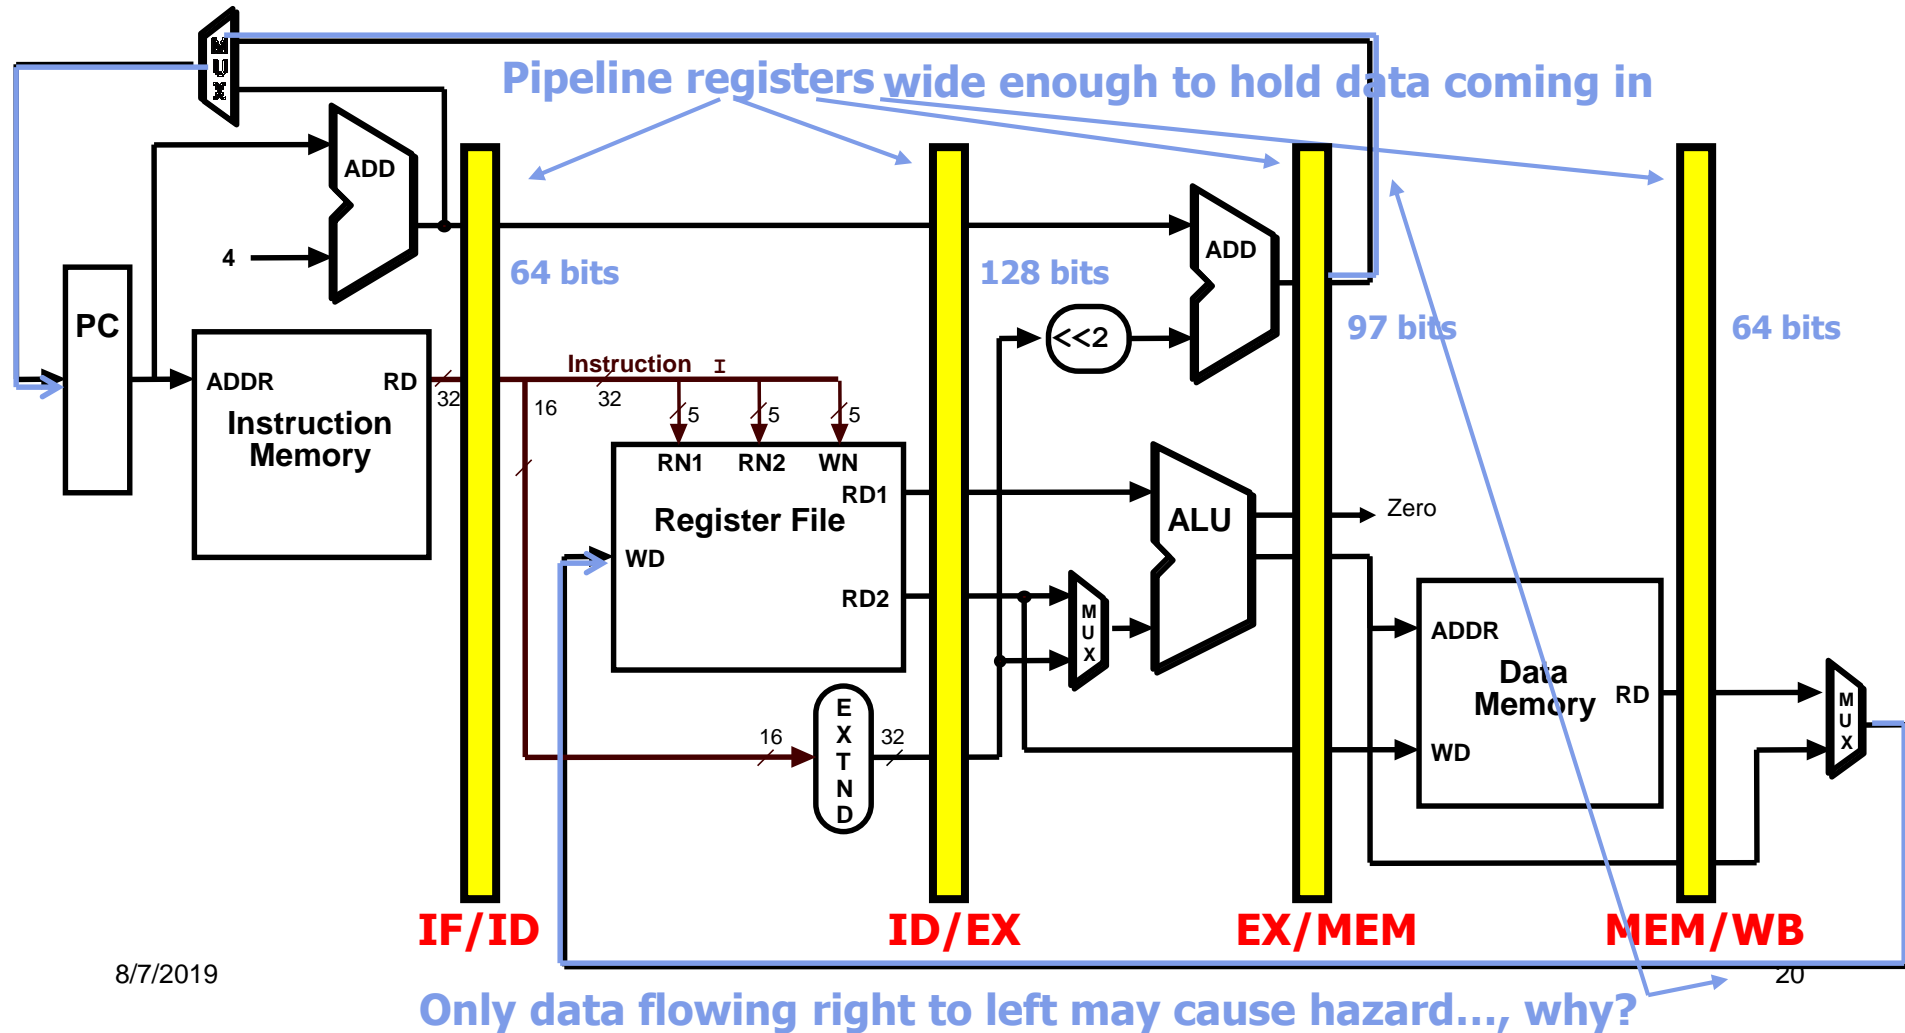


- *What happens if we break the execution into multiple cycles, but keep the extra hardware?*
 - Answer: *We may be able to start executing a new instruction at each clock cycle - pipelining*
- ...but we shall need *extra* registers to hold data between cycles – *pipeline registers*

Pipelined Datapath

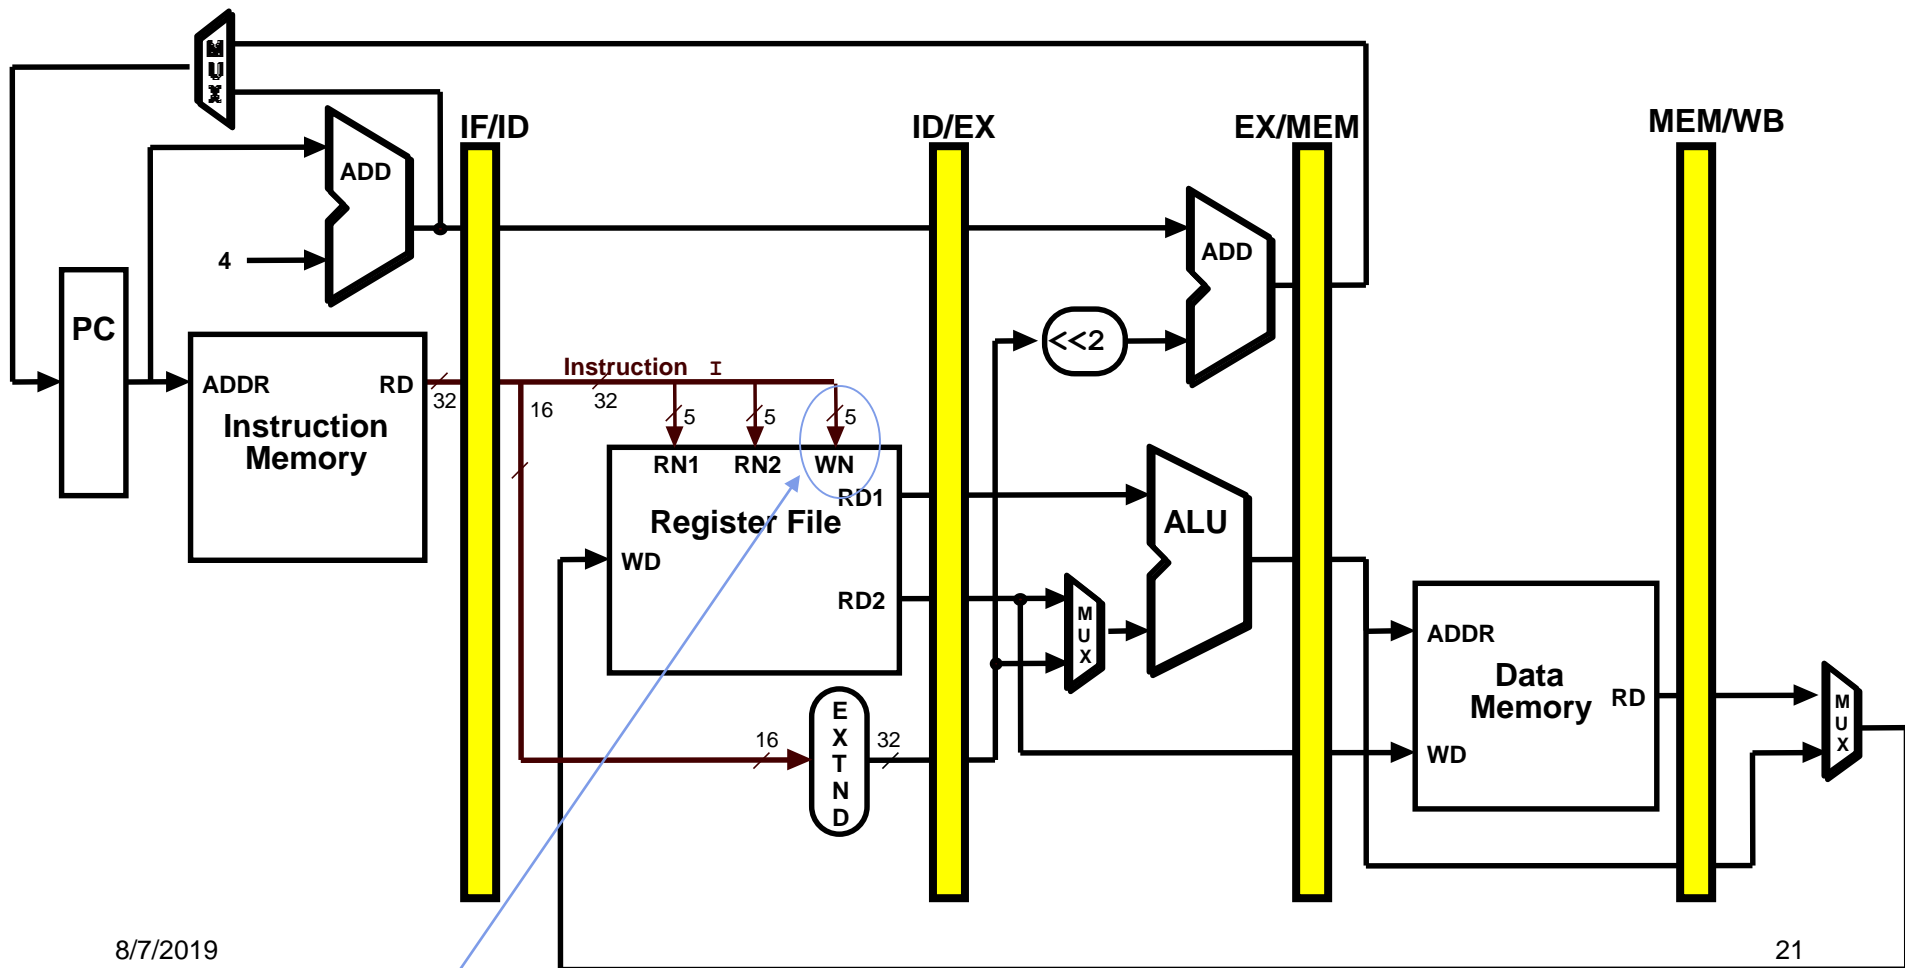


Pipelined Datapath





Bug in the Datapath

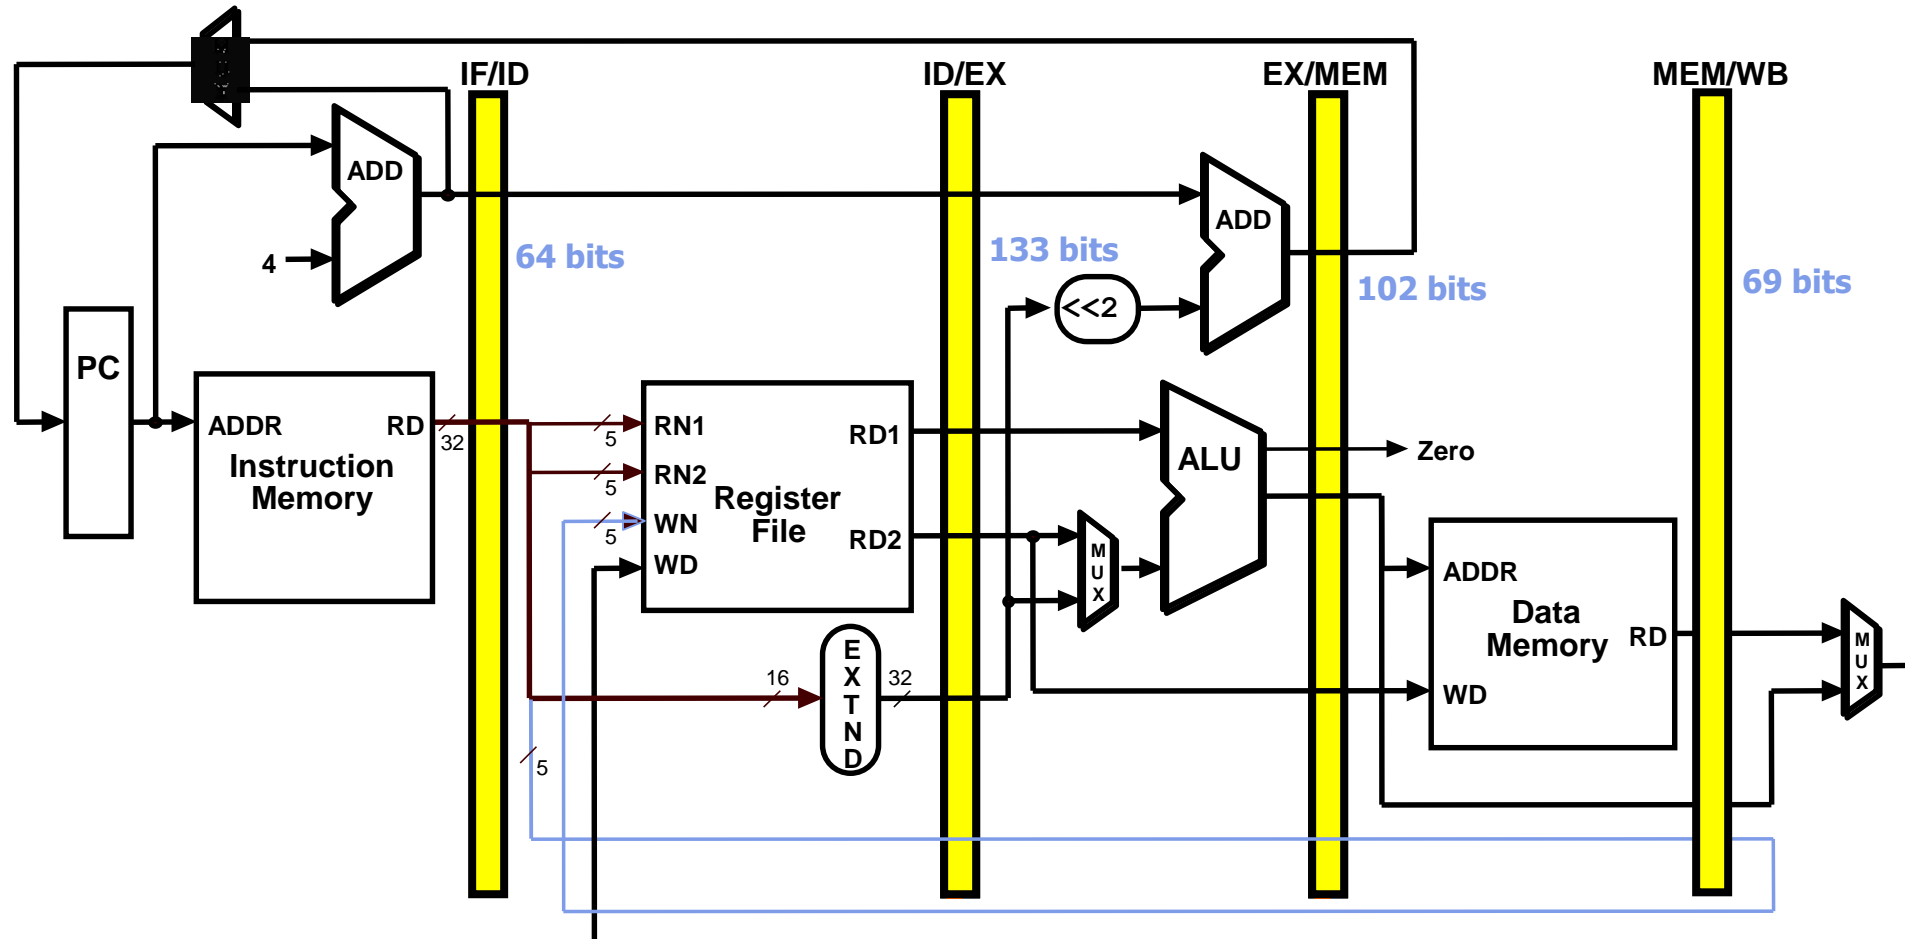


8/7/2019

Write register number comes from another *later* instruction!

21

Corrected Datapath





Pipelined Example

- Consider the following instruction sequence:

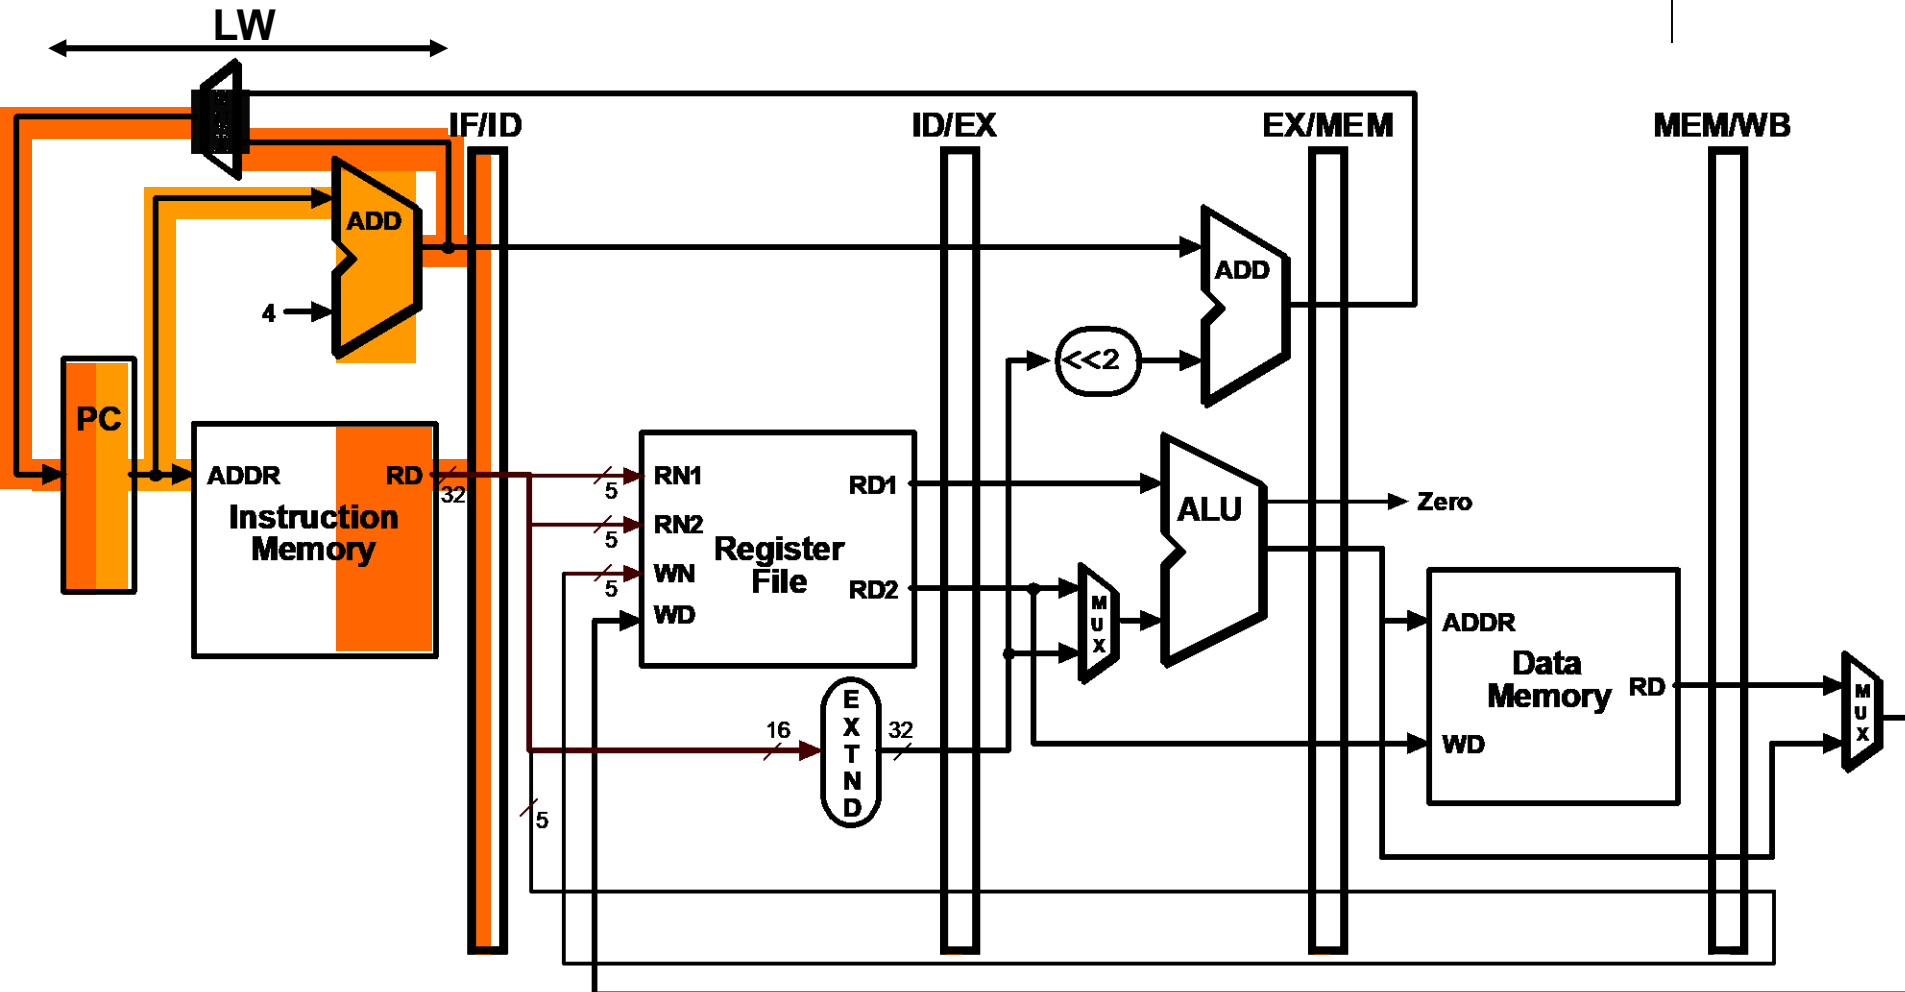
```
lw    $t0, 10($t1)
```

```
sw    $t3, 20($t4)
```

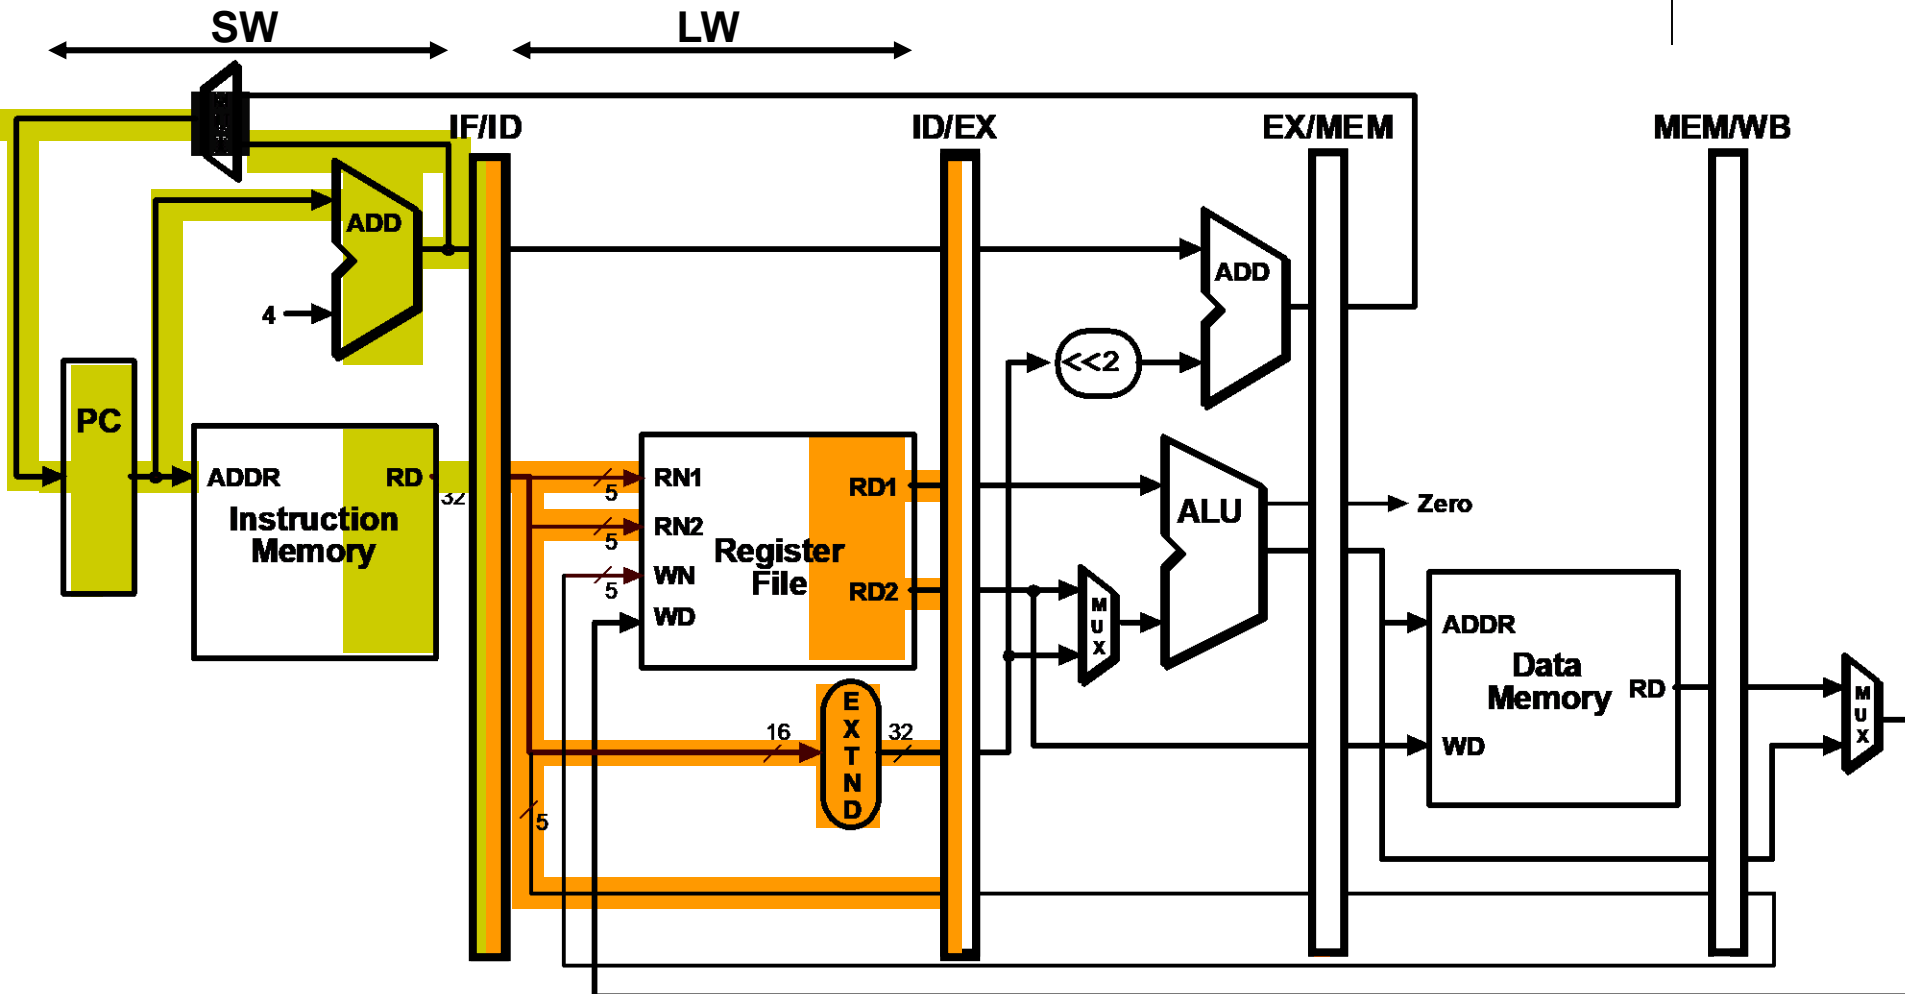

```
add   $t5, $t6, $t7
```

```
sub   $t8, $t9, $t10
```

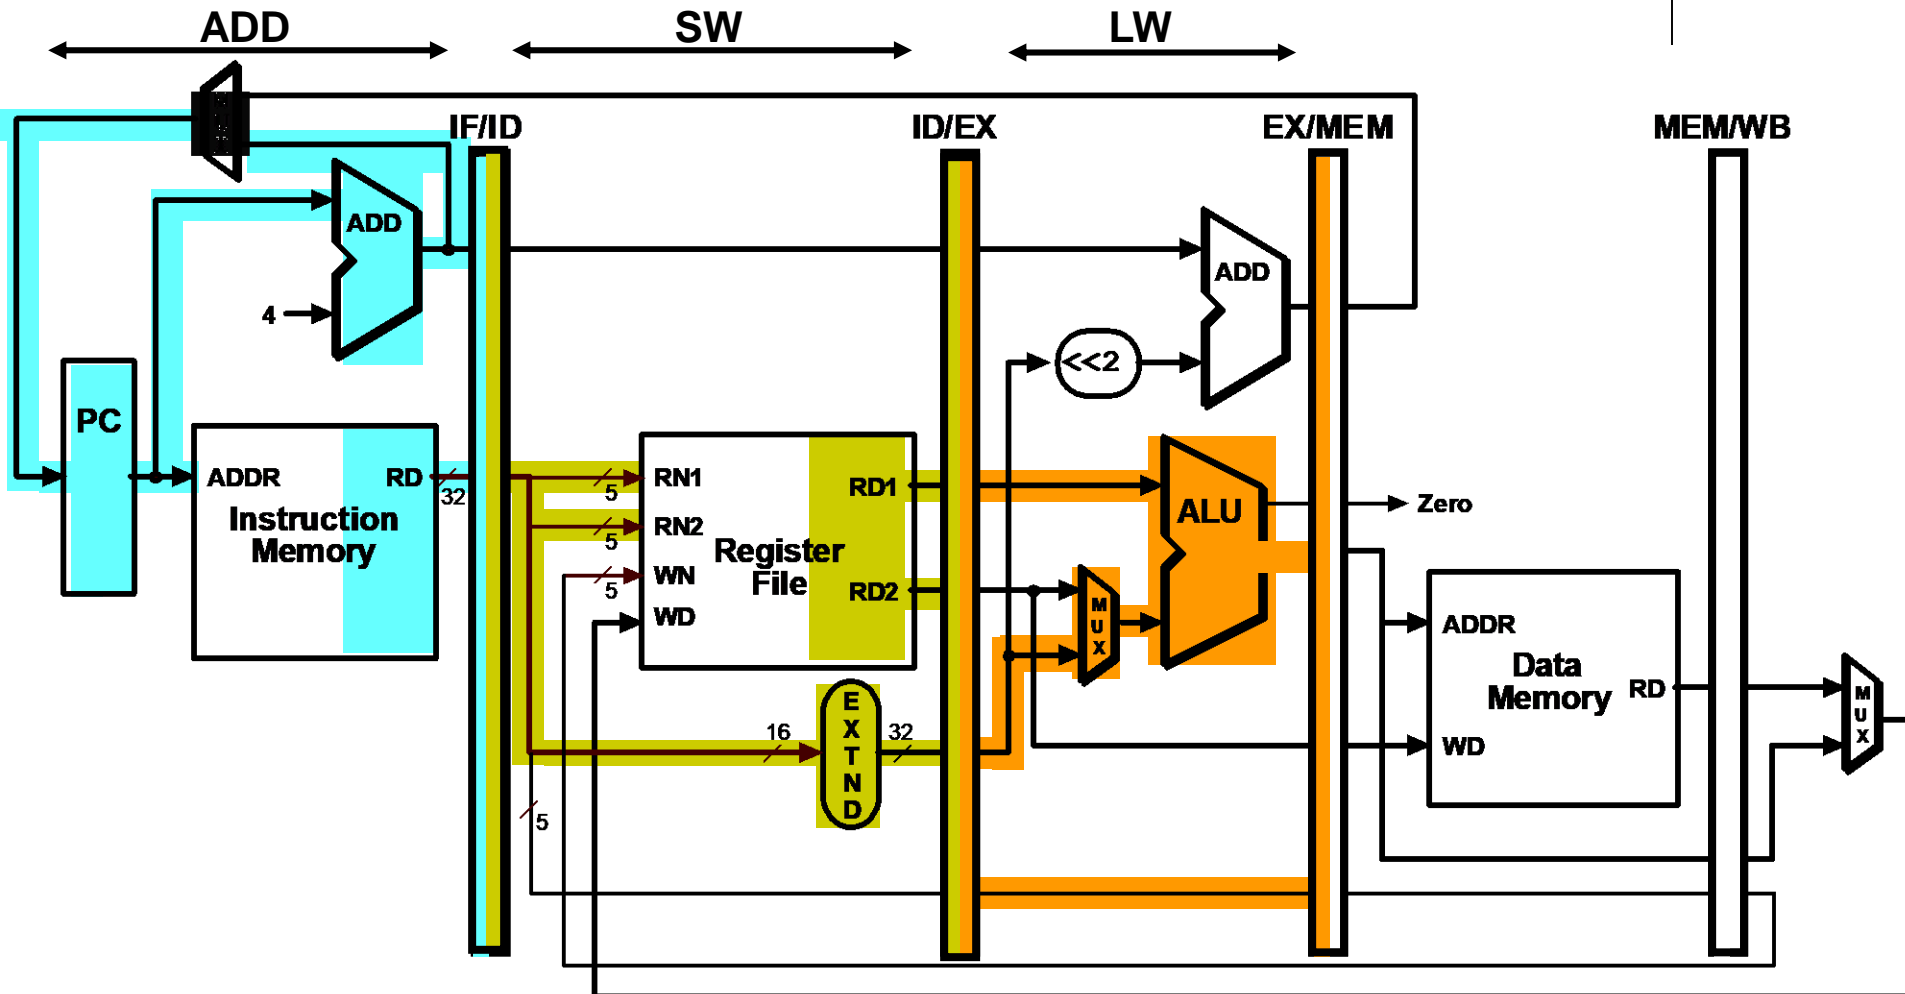

Single-Clock-Cycle Diagram: Clock Cycle 1



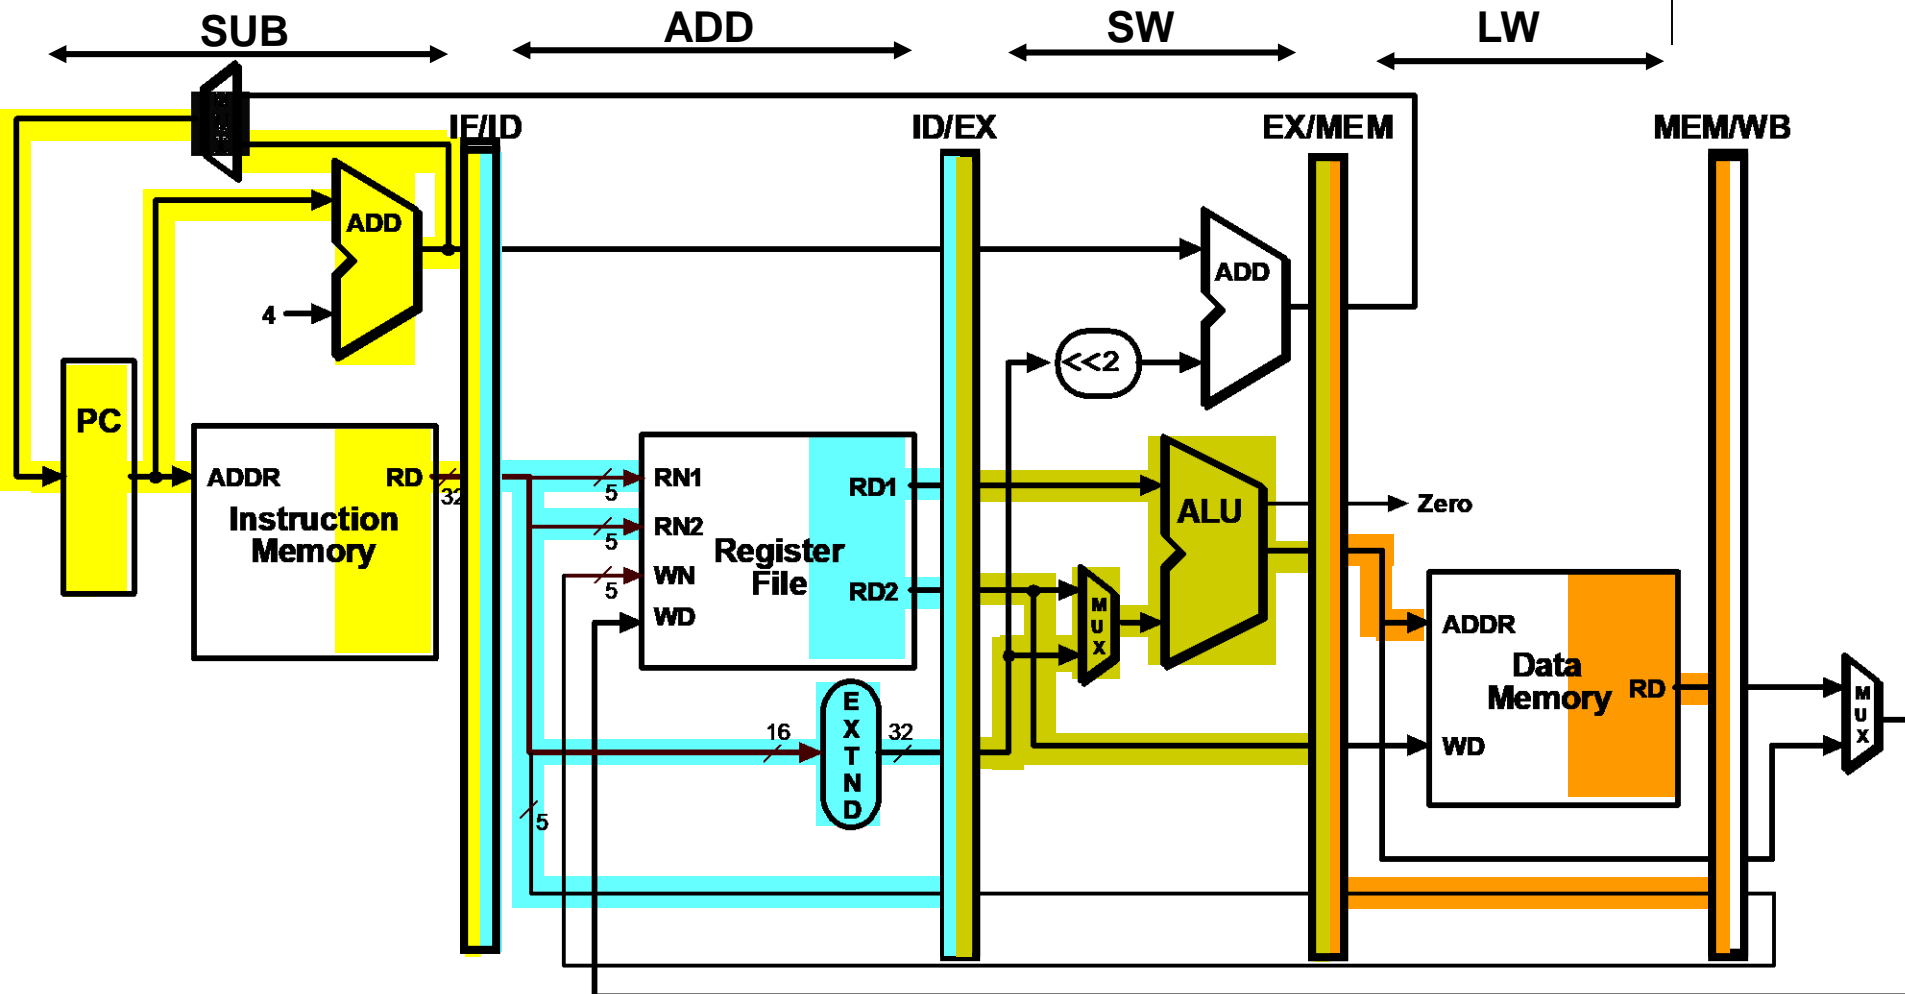

Single-Clock-Cycle Diagram: Clock Cycle 2



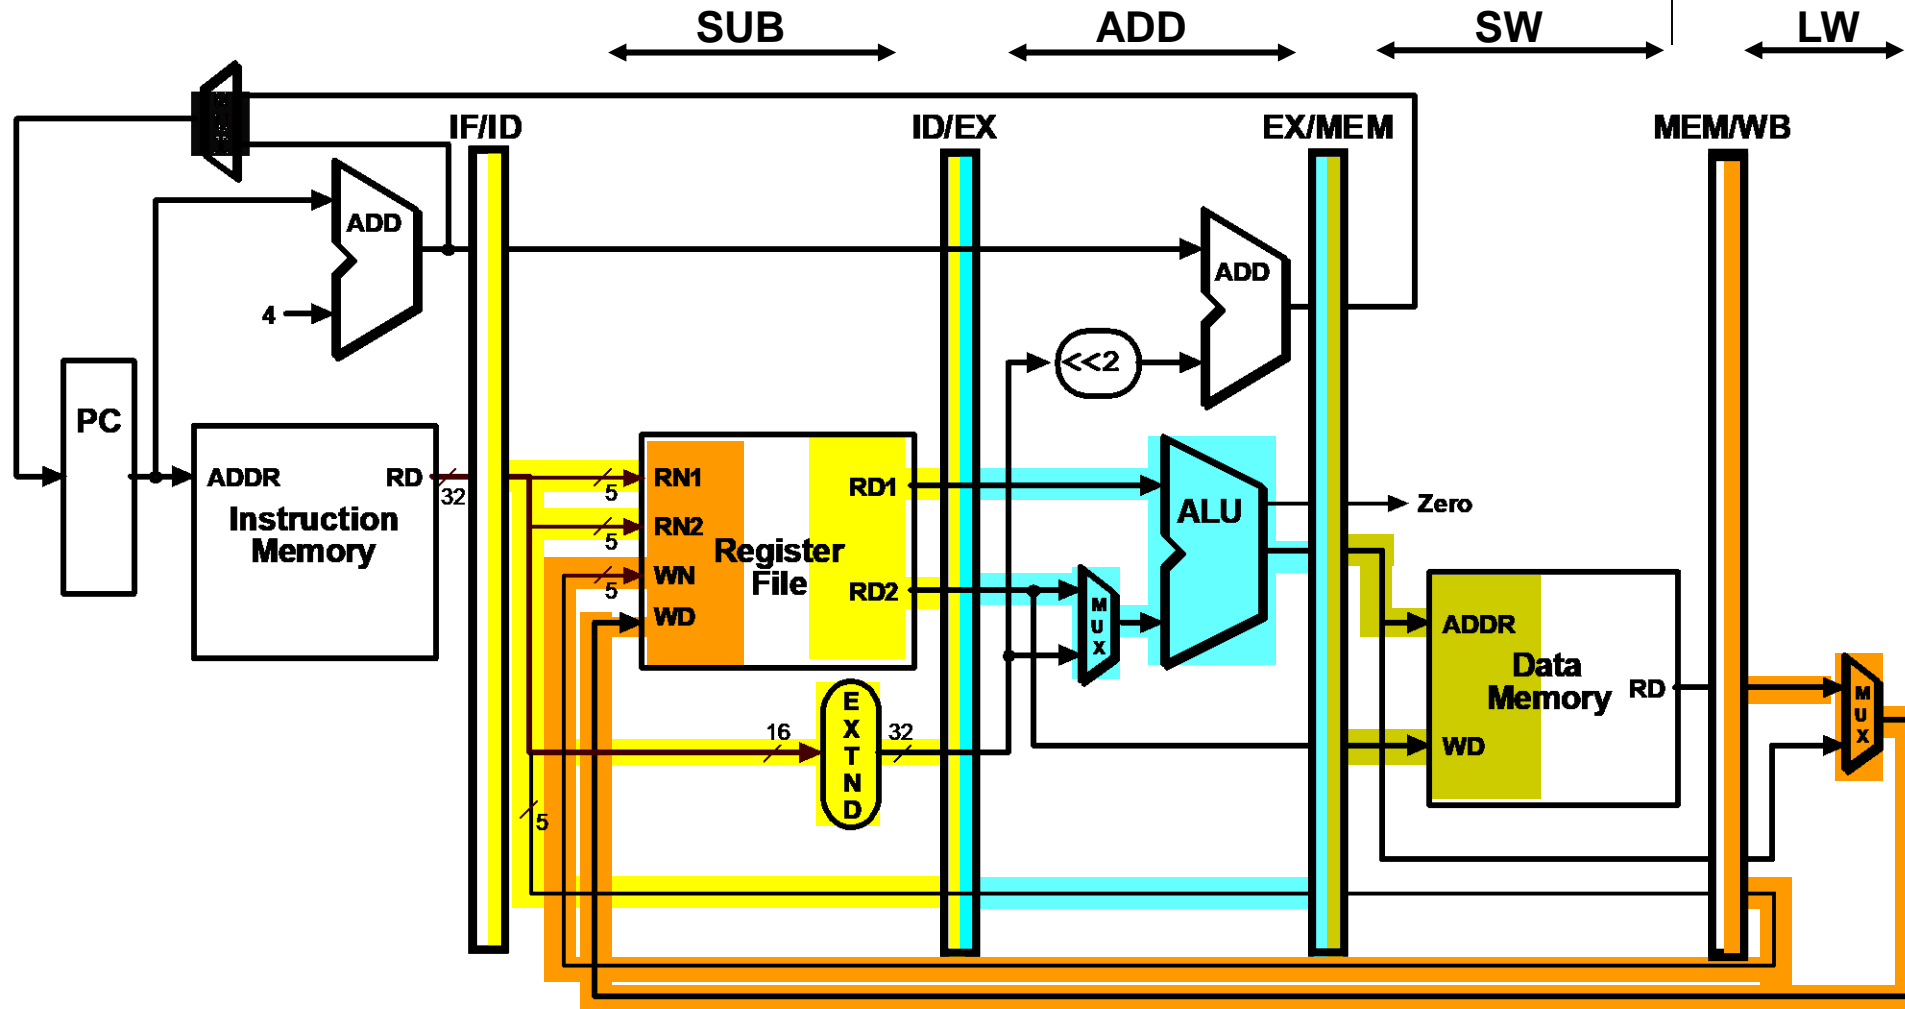
Single-Clock-Cycle Diagram: Clock Cycle 3



Single-Clock-Cycle Diagram: Clock Cycle 4

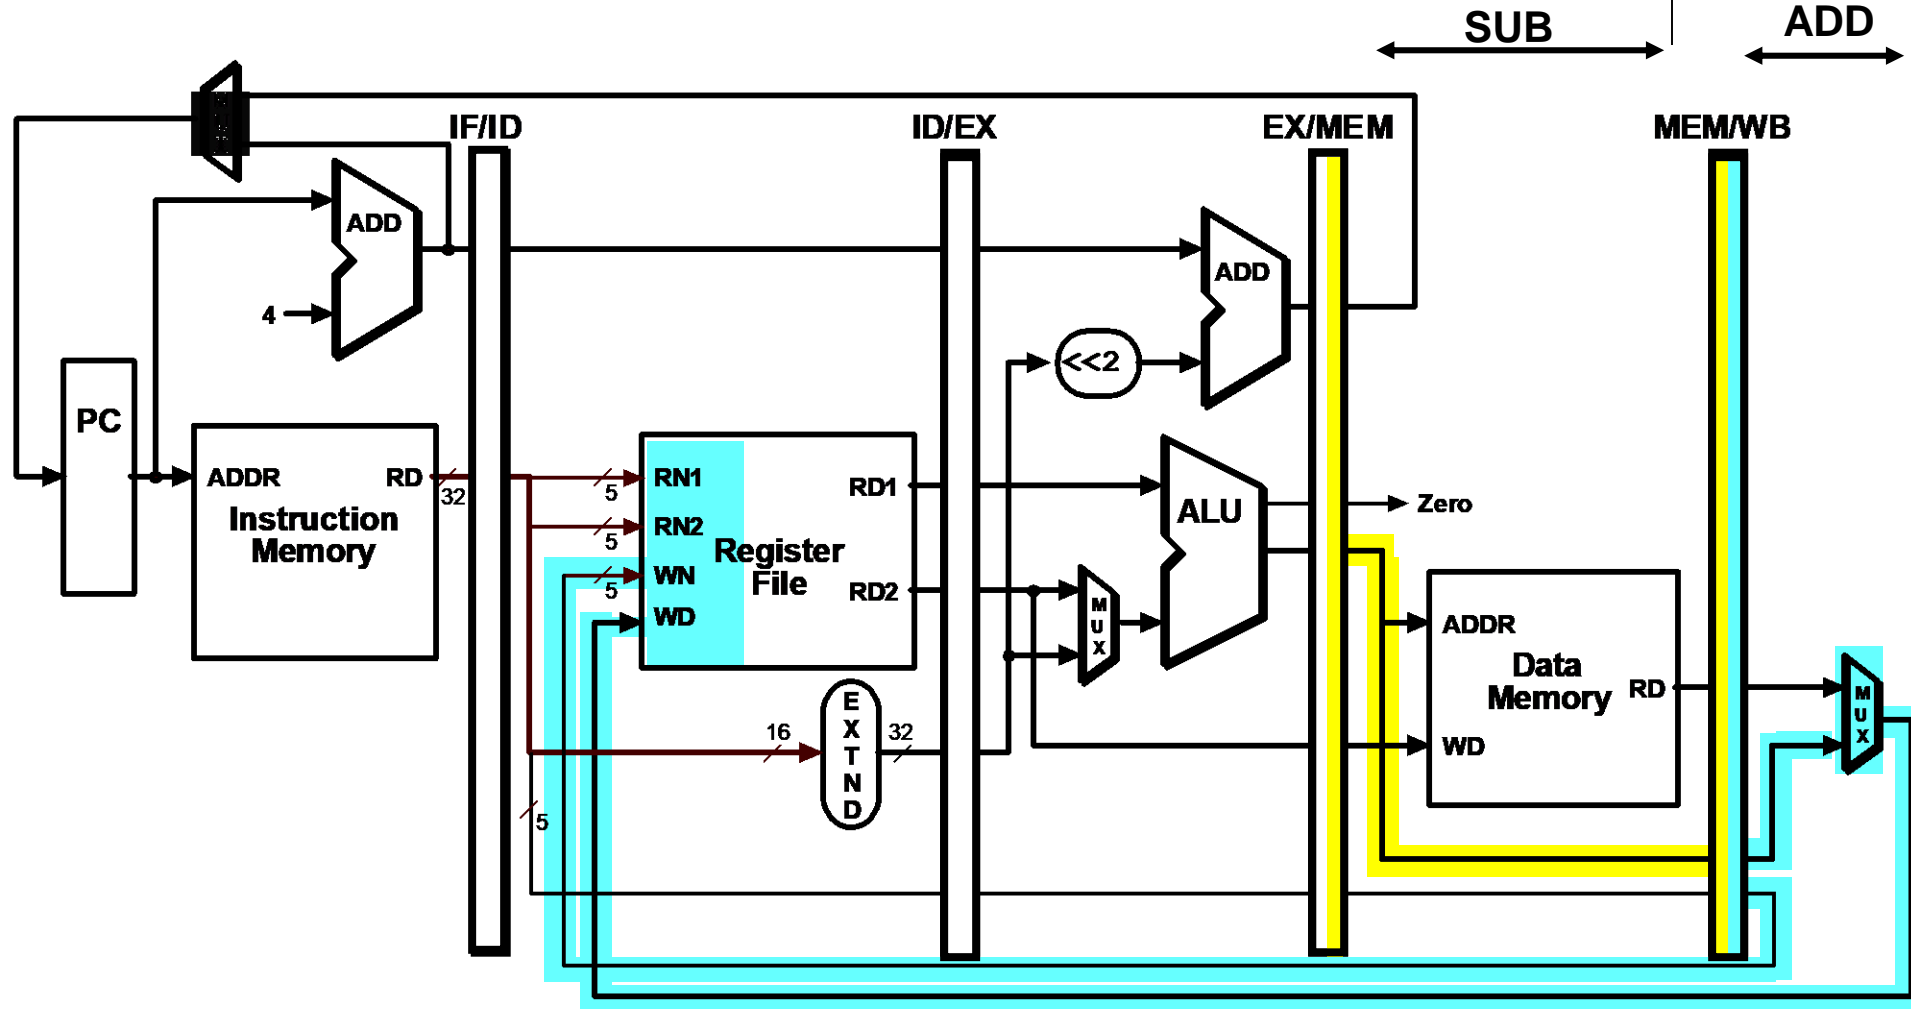



Single-Clock-Cycle Diagram: Clock Cycle 5

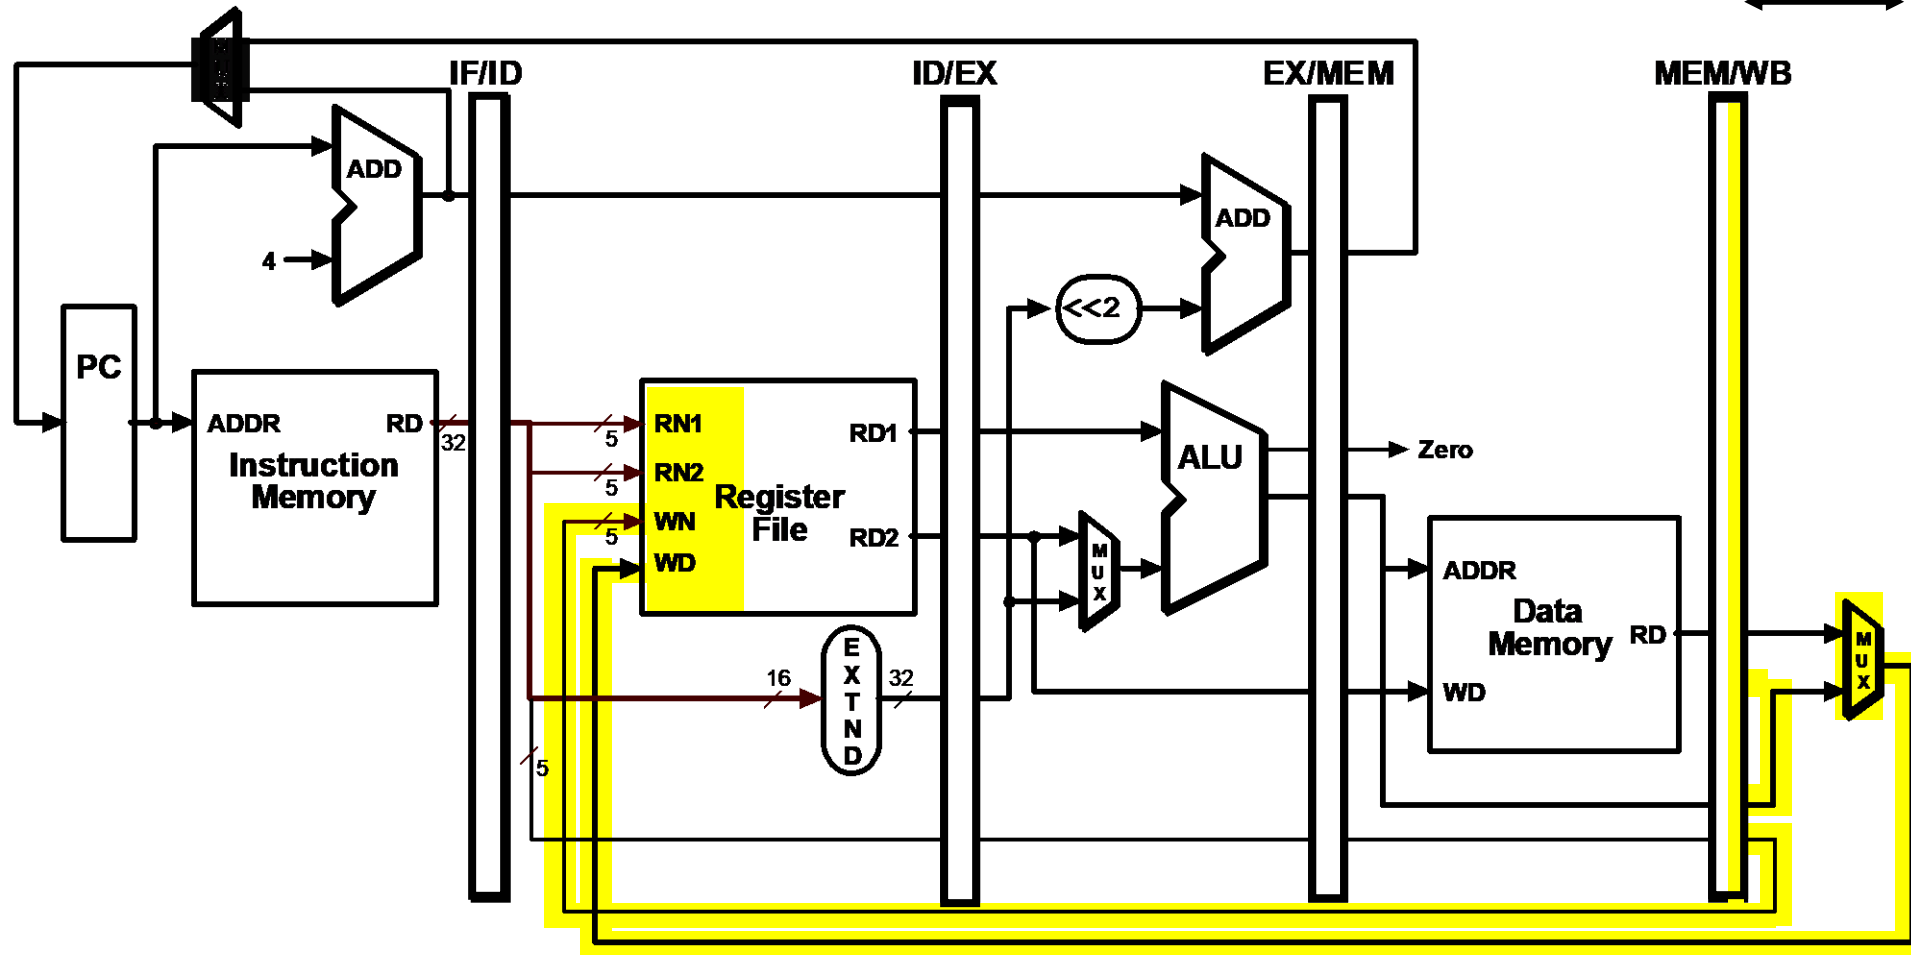





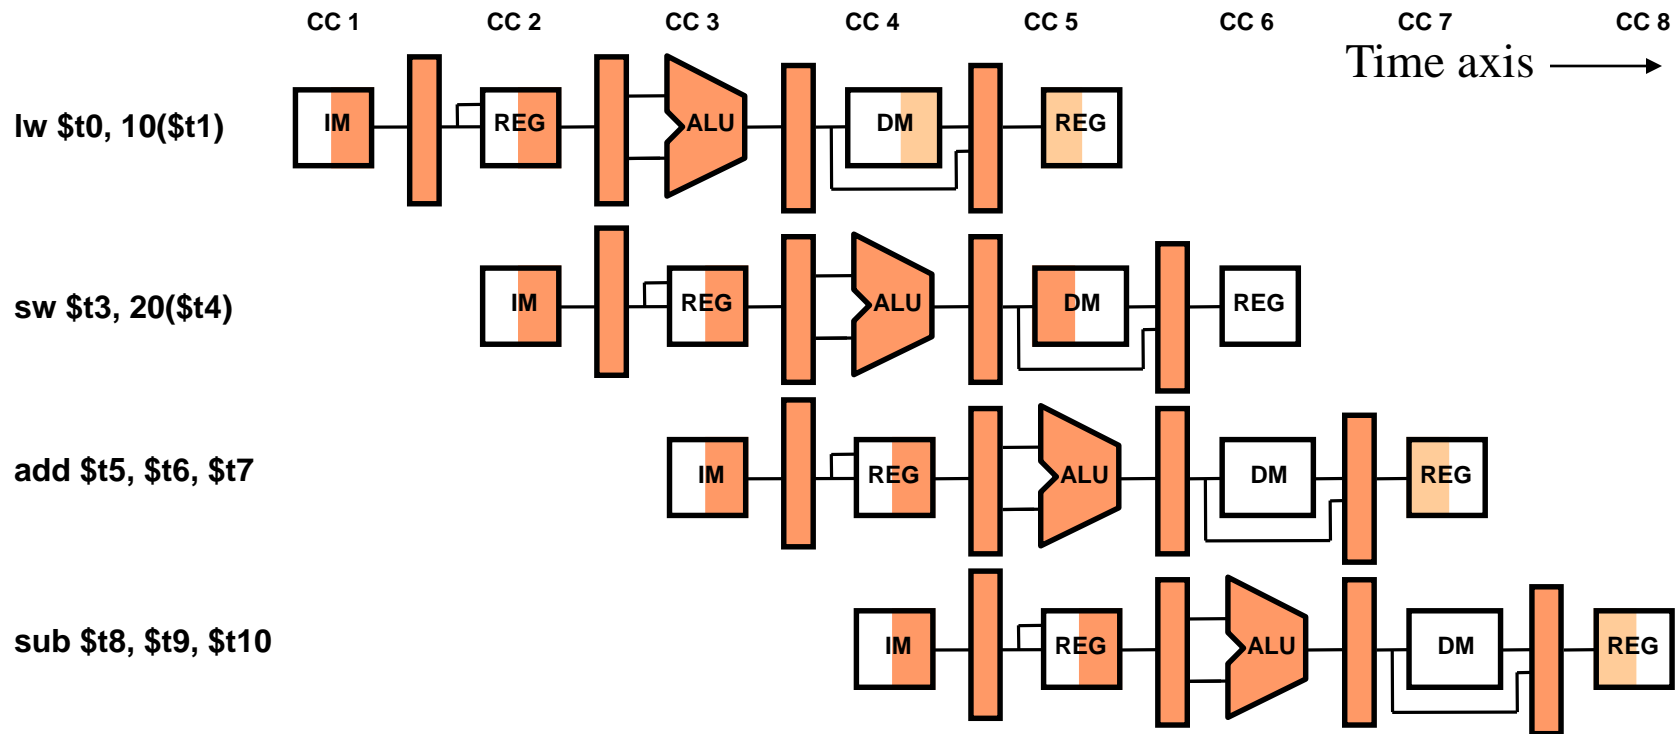
Single-Clock-Cycle Diagram: Clock Cycle 7



Single-Clock-Cycle Diagram: Clock Cycle 8



Alternative View – Multiple-Clock-Cycle Diagram





Notes

- One significant difference in the execution of an R-type instruction between multicycle and pipelined implementations:
 - register write-back for the R-type instruction is the 5th (the last write-back) pipeline stage vs. the 4th stage for the multicycle implementation. *Why?*
 - think of *structural hazards* when writing to the register file...
- Worth repeating: the *essential difference* between the pipeline and multicycle implementations is the insertion of pipeline registers to *decouple the 5 stages*
- The CPI of an *ideal pipeline* (no stalls) is 1. *Why?*

Simple Example: Comparing Performance



- *Compare performance for multicycle, and pipelined datapaths using the gcc instruction mix*
 - assume 2 ns for memory access
 - assume gcc instruction mix 23% loads, 13% stores, 19% branches, 2% jumps, 43% ALU
 - for pipelined execution assume
 - 50% of the loads are followed immediately by an instruction that uses the result of the load. This sacrifices 2 cycles.
 - 25% of branches are mispredicted
 - branch delay on misprediction is 1 clock cycle
 - jumps always incur 1 clock cycle delay so their average time is 2 clock cycles

Simple Example: Comparing Performance



- *Multicycle*: average instruction time 8.04 ns
- *Pipelined*:
 - loads use 1 cc (clock cycle) when no load-use dependency and 2 cc when there is dependency – given 50% of loads are followed by dependency the average cc per load is:
 $0.5 \times 1 + 0.5 \times 2 = 1.5$
 - stores use 1 cc each
 - branches use 1 cc when predicted correctly and 2 cc when not – given 25% mis-prediction average cc per branch is
 $0.75 \times 1 + 0.25 \times 2 = 1.25$
 - jumps use 3 cc each
 - ALU instructions use 1 cc each
 - therefore, average CPI is
 $1.5 \times 23\% + 1 \times 13\% + 1.25 \times 19\% + 3 \times 2\% + 1 \times 43\% = 1.1825$
 - therefore, average instruction time is $1.1825 \times 2 = 2.365$ ns

Summary



- *Techniques described in this chapter to design datapaths and control are at the core of all modern computer architecture*
- Multicycle datapaths offer two great advantages over single-cycle
 - functional units can be reused within a single instruction if they are accessed in different cycles – reducing the need to replicate expensive logic
 - instructions with shorter execution paths can complete quicker by consuming fewer cycles
- Modern computers, in fact, take the multicycle paradigm to a higher level to achieve greater instruction throughput:
 - *pipelining* (next topic) where multiple instructions execute simultaneously by having cycles of different instructions overlap in the datapath
 - *the MIPS architecture was designed to be pipelined*